

AI models keep getting caught cheating

Report Type: Weekly · Generated: July 21, 2026 21:26 UTC · Coverage: Jul 21 - Jul 21, 2026

1 Total Advisories	0 Critical	0 High	0 Medium
LOW Severity	N/A CVSS / Risk Score	TLP:GREEN TLP Classification	PRO Customer Tier

Executive Summary

[P4] Phishing - low severity (Risk 1.6/10). MITRE ATT&CK mapped: T1059, T1190 (2 techniques). Actor attribution: CDB-UNATTR-APT. AI confidence: LOW (23%) - limited source data available.

CONFIDENTIALITY NOTICE: This report is classified PRO tier and intended exclusively for the named recipient. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited. © 2026 CYBERDUDEBIVASH SENTINEL APEX. All rights reserved.

Top Threat Advisories

High-priority advisories ranked by CVSS score, exploitability, and actor attribution.

Advisory Title	Severity	CVSS	EPSS	AI Summary
AI models keep getting caught cheating	LOW	-	-	[P4] Phishing - low severity (Risk 1.6/10). MITRE ATT&CK mapped to T1059, T1190.

MITRE ATT&CK® v15 Mapping

Technique density score: 0.4 · Techniques observed: 2 · Tactics covered: 2

Technique ID	Name	Tactic	Count
T1059	Command and Scripting Interpreter	Execution	1
T1190	Exploit Public-Facing Application	Initial Access	1

Tactics Covered

Execution	Initial Access
1 technique(s)	1 technique(s)

Threat Actor Intelligence

Attribution analysis across advisories in this report period.

Threat Actor	Count	Associated CVEs / Indicators
UNC-CDB	1	-

Indicators of Compromise (IOCs)

IOC extraction for this advisory is pending enrichment pipeline completion. PRO subscribers receive real-time IOC push to SIEM/SOAR via the SENTINEL APEX webhook at: **GET /api/v1/intel/iocs?advisory={id}**. Full IOC packages include IPv4, domain, SHA256, and URL indicators with confidence scores, first-seen timestamps, and STIX 2.1 formatted bundles.

Detection Engineering

Deploy these production-ready detection rules into your SIEM platform to identify exploitation attempts in real time. Compatible with Microsoft Sentinel, Splunk ES, and any Sigma-compatible platform.

Sigma Rule - Webserver / Process Monitoring

```
title: SENTINEL APEX - CDB-ADVISORY Exploitation Attempt
id: cdb-cdb-advisory-det
status: experimental
description: Detects exploitation indicators related to CDB-ADVISORY.
references:
- https://intel.cyberdudebivash.com
author: CyberDudeBivash SENTINEL APEX
date: 2026/07/21
tags:
- attack.initial_access
- attack.t1190
logsource:
category: webserver
detection:
keywords:
- 'CDB-ADVISORY'
condition: keywords
falsepositives:
- Security testing / penetration testing
level: medium
```

Microsoft Sentinel - KQL Query

```
// SENTINEL APEX - AI models keep getting caught cheating
// Advisory | Microsoft Sentinel KQL
let time_window = ago(24h);
SecurityAlert
| where TimeGenerated > time_window
| where Description has "AI models keep getting caught cheating"
or AlertName has "AI models keep getting caught cheating"
or ExtendedProperties has "AI models keep getting caught cheating"
| extend Rule = "APEX-ADVISORY", Severity = "LOW"
| project TimeGenerated, AlertName, AlertSeverity, Entities, Description, Rule, Severity
| order by TimeGenerated desc
```

Incident Response Playbook

Phase 1 - 0-24 Hours (Containment)

1. Activate IR team; assign lead analyst and executive sponsor.
2. Isolate or WAF-shield AI models keep getting caught cheating instances exposed to the internet.
3. Enable verbose access logging on all AI models keep getting caught cheating endpoints immediately.
4. Pull last 72h of web server and application logs for forensic baselining.

- 5. Check SENTINEL APEX IOC feeds and SIEM alerts for exploitation hits in your environment.
- 6. Notify relevant internal stakeholders and legal / compliance team.

Phase 2 - 24-72 Hours (Eradication)

- 1. Apply vendor patch for AI models keep getting caught cheating or implement recommended mitigation.
- 2. Deploy Sigma and KQL detection rules from Section 6 across all SIEM platforms.
- 3. Rotate all API keys, service account credentials, and session tokens.
- 4. Conduct forensic timeline reconstruction for potential compromise window.
- 5. Validate patch deployment across 100% of affected instances via asset inventory.

Phase 3 - 7 Days (Recovery & Hardening)

- 1. Conduct post-incident review and update incident response runbooks.
- 2. Threat hunt for lateral movement or persistence artifacts using MITRE ATT&CK mapping.
- 3. Update asset inventory with patched version information and closure evidence.
- 4. Submit findings to internal risk register and CISO dashboard.
- 5. Schedule follow-up vulnerability scan to confirm remediation completeness.

Financial Impact Analysis (FAIR Model)

\$20K - \$180K Breach Cost Range (FAIR)	\$60K IBM Cost of Breach 2025 Median	\$25K/day Business Interruption Cost	N/A Regulatory Fine Exposure
---	--	--	--

Low severity findings require targeted patches. Aggregated low-severity debt compounds to medium breach probability within 12 months.

Remediation Cost Estimate: \$15K · Includes engineering time, IR retainer activation, forensic investigation, and customer notification costs.

Regulatory & Compliance Implications

Organizations processing this advisory must evaluate obligations under applicable regulatory frameworks. The table below maps this threat to relevant compliance requirements.

Framework	Reference	Obligation
ISO 27001:2022	A.8.8	Track and remediate within standard 90-day patch cycle
NIST CSF 2.0	ID.RA	Include in regular vulnerability risk assessment

Actionable Recommendations

Prioritized defensive actions based on this period's intelligence.

1. Monitor AI models keep getting caught cheating for escalation.
2. Apply patches in next maintenance window.
3. Apply vendor patch for AI models keep getting caught cheating immediately - check NVD for official fix guidance.
4. Enable enhanced logging and alerting on all systems running affected software versions.
5. Deploy the Sigma detection rule from this report into your SIEM platform (Sentinel / Splunk / Elastic).
6. Audit all internet-facing instances of affected products and confirm patch deployment.
7. Conduct a threat hunt using MITRE ATT&CK techniques mapped in this advisory.

Appendix - Report Metadata & Legal

Field	Value
Report ID	intel--2ee927892b6f14a293ffc12b
Report Type	Weekly
Platform	CYBERDUDEBIVASH SENTINEL APEX v166.0 GOD-MODE
Generated At	2026-07-21T21:26:13.50107
Customer Tier	PRO
Customer Email	-
TLP	TLP:GREEN
STIX Version	2.1
ATT&CK Version	v15

Legal & Usage Notice

This threat intelligence report is produced by CYBERDUDEBIVASH SENTINEL APEX, operated by Bivash Nath (GSTIN: 21ARKPN8270G1ZP). All intelligence is aggregated from public vulnerability databases (NVD, CVSS), threat feeds, and proprietary AI enrichment pipelines. This report is provided AS-IS for informational and defensive security purposes only. CYBERDUDEBIVASH accepts no liability for actions taken based on this intelligence. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited without written consent. For licensing inquiries: enterprise@cyberdudebivash.com

CYBERDUDEBIVASH® SENTINEL APEX // CONFIDENTIAL - FOR AUTHORIZED USE ONLY