

gemini-mcp-tool vulnerable to OS command injection and @file exfiltration via prompt quoting (CVE-2026-0755)

Report Type: Weekly · Generated: June 18, 2026 22:39 UTC · Coverage: Jun 18 - Jun 18, 2026

1 Total Advisories	1 Critical	0 High	0 Medium
CRITICAL Severity	9.8 CVSS / Risk Score	TLP:CLEAR TLP Classification	PRO Customer Tier

Executive Summary

Untrusted prompt input could reach the Gemini CLI @file parser, allowing read/exfiltration of arbitrary local files (@/etc/passwd, @~/.ssh/id_rsa, @../secret). On Windows, unquoted cmd.exe metacharacters could break out into OS command injection. Fix (1.1.6): removed the broken shell:false double-quote wrapping; added assertSafeFileReferences() to contain @file refs to the working directory; hardened Windows cmd.exe argument quoting.

CONFIDENTIALITY NOTICE: This report is classified PRO tier and intended exclusively for the named recipient. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited. © 2026 CYBERDUDEBIVASH SENTINEL APEX. All rights reserved.

CYBERDUDEBIVASH® SENTINEL APEX // CONFIDENTIAL - FOR AUTHORIZED USE ONLY

Top Threat Advisories

High-priority advisories ranked by CVSS score, exploitability, and actor attribution.

Advisory Title	Severity	CVSS	EPSS	AI Summary
gemini-mcp-tool vulnerable to OS command injection and file exfiltration	CRITICAL	9.8	3.3%	Untrusted prompt input could reach the Gemini CLI @file parser, a

Vulnerability Intelligence Deep-Dive

9.8 CVSS 3.1 Score	3.3% EPSS 30-day Probability	NO CISA KEV Listed	CRITICAL Severity Rating
--	--	--	--

Field	Value
CVE / Advisory ID	CVE-2026-0755
Vulnerability Class	Remote Code Execution (CWE-78/CWE-94)
CWE Reference	Pending NVD enrichment
Actor Attribution	CDB-UNATTR-CVE
Source Advisory	NVD / Vendor Advisory

Refer to the official NVD entry and vendor advisory for the complete list of affected products and versions for CVE-2026-0755.

Vulnerability Context & Attack Surface

Remote code execution vulnerabilities allow attackers to execute arbitrary operating system commands in the context of the application server process. Successful exploitation typically results in full server compromise, lateral movement, and persistent backdoor installation. Internet-facing applications are prioritized for exploitation - median time from PoC publication to in-the-wild exploitation is 3-14 days.

MITRE ATT&CK® v15 Mapping

Technique density score: 0.0 · Techniques observed: 0 · Tactics covered: 0

MITRE ATT&CK techniques pending analyst enrichment. Deploy detection rules from Section 6 to identify related activity.

Threat Actor Intelligence

Attribution analysis across advisories in this report period.

Threat Actor	Count	Associated CVEs / Indicators
CDB-UNATTR-CVE	1	-

Indicators of Compromise (IOCs)

IOC extraction for this advisory is pending enrichment pipeline completion. PRO subscribers receive real-time IOC push to SIEM/SOAR via the SENTINEL APEX webhook at: GET /api/v1/intell/iocs?advisory={id}. Full IOC packages include IPv4, domain, SHA256, and URL indicators with confidence scores, first-seen timestamps, and STIX 2.1 formatted bundles.

Detection Engineering

Deploy these production-ready detection rules into your SIEM platform to identify exploitation attempts in real time. Compatible with Microsoft Sentinel, Splunk ES, and any Sigma-compatible platform.

Sigma Rule - Webserver / Process Monitoring

```
title: SENTINEL APEX - CVE-2026-0755 Exploitation Attempt
id: cdb-cve-2026-0755-det
status: experimental
description: Detects exploitation indicators related to CVE-2026-0755.
references:
- https://intel.cyberdudebivash.com
author: CyberDudeBivash SENTINEL APEX
date: 2026/06/18
tags:
- attack.initial_access
- attack.t1190
logsource:
category: webserver
detection:
keywords:
- 'CVE-2026-0755'
condition: keywords
falsepositives:
- Security testing / penetration testing
level: medium
```

Microsoft Sentinel - KQL Query

```
// SENTINEL APEX - gemini-mcp-tool vulnerable to OS command injection and @file exfiltration via prompt quoting (CVE-2026-0755)
// CVE-2026-0755 | Microsoft Sentinel KQL
let time_window = ago(24h);
SecurityAlert
| where TimeGenerated > time_window
| where Description has "CVE-2026-0755"
or AlertName has "CVE-2026-0755"
or ExtendedProperties has "CVE-2026-0755"
| extend Rule = "APEX-CVE20260755", Severity = "CRITICAL"
| project TimeGenerated, AlertName, AlertSeverity, Entities, Description, Rule, Severity
| order by TimeGenerated desc
```

Incident Response Playbook

Phase 1 - 0-24 Hours (Containment)

1. Activate IR team; assign lead analyst and executive sponsor.
2. Isolate or WAF-shield gemini instances exposed to the internet.
3. Enable verbose access logging on all gemini endpoints immediately.

- 4. Pull last 72h of web server and application logs for forensic baselining.
- 5. Check SENTINEL APEX IOC feeds and SIEM alerts for exploitation hits in your environment.
- 6. Notify relevant internal stakeholders and legal / compliance team.

Phase 2 - 24-72 Hours (Eradication)

- 1. Apply vendor patch for CVE-2026-0755 or implement recommended mitigation.
- 2. Deploy Sigma and KQL detection rules from Section 6 across all SIEM platforms.
- 3. Rotate all API keys, service account credentials, and session tokens.
- 4. Conduct forensic timeline reconstruction for potential compromise window.
- 5. Validate patch deployment across 100% of affected instances via asset inventory.

Phase 3 - 7 Days (Recovery & Hardening)

- 1. Conduct post-incident review and update incident response runbooks.
- 2. Threat hunt for lateral movement or persistence artifacts using MITRE ATT&CK mapping.
- 3. Update asset inventory with patched version information and closure evidence.
- 4. Submit findings to internal risk register and CISO dashboard.
- 5. Schedule follow-up vulnerability scan to confirm remediation completeness.

Financial Impact Analysis (FAIR Model)

\$2.4M - \$12M Breach Cost Range (FAIR)	\$5.8M IBM Cost of Breach 2025 Median	\$1.2M/day Business Interruption Cost	\$850K Regulatory Fine Exposure
---	---	---	---

Critical severity breaches trigger mandatory 72-hour regulatory disclosure (GDPR Art. 33), class-action exposure, and long-tail legal costs averaging 18-36 months post-incident.

Remediation Cost Estimate: \$380K · Includes engineering time, IR retainer activation, forensic investigation, and customer notification costs.

Regulatory & Compliance Implications

Organizations processing this advisory must evaluate obligations under applicable regulatory frameworks. The table below maps this threat to relevant compliance requirements.

Framework	Reference	Obligation
GDPR	Art. 33 / Art. 83(4)	72-hour breach notification to DPA; fines up to 4% of global annual turnover
PCI-DSS v4.0	Req. 6.3 / 6.4	Critical patches within 1 month; compensating controls required
NIS2 Directive	Art. 23	Early warning within 24h; full incident report within 72h

Framework	Reference	Obligation
SOX / ICFR	Section 302 / 404	Material cybersecurity weakness disclosure in 10-K/10-Q
ISO 27001:2022	A.8.8 / A.5.29	Documented vulnerability management and incident management

Actionable Recommendations

Prioritized defensive actions based on this period's intelligence.

1. Apply patches immediately for gemini-mcp-tool vulnerable to OS command injection and @file exfiltration via prompt quoting (CVE-2026-0755).
2. Enable enhanced monitoring for related IOCs.
3. Review SIEM rules for associated MITRE ATT&CK techniques.
4. Apply vendor patch for gemini-mcp-tool vulnerable to OS command injection and @file exfiltration via prompt quoting (CVE-2026-0755) immediately - check NVD for official fix guidance.
5. Enable enhanced logging and alerting on all systems running affected software versions.
6. Deploy the Sigma detection rule from this report into your SIEM platform (Sentinel / Splunk / Elastic).
7. Audit all internet-facing instances of affected products and confirm patch deployment.

Appendix - Report Metadata & Legal

Field	Value
Report ID	intel--841191f6fe93cd9589
Report Type	Weekly
Platform	CYBERDUDEBIVASH SENTINEL APEX v166.0 GOD-MODE
Generated At	2026-06-18T22:39:23.46465
Customer Tier	PRO
Customer Email	-
TLP	TLP: CLEAR
STIX Version	2.1
ATT&CK Version	v15

Legal & Usage Notice

This threat intelligence report is produced by CYBERDUDEBIVASH SENTINEL APEX, operated by Bivash Nath (GSTIN: 21ARKPN8270G1ZP). All intelligence is aggregated from public vulnerability databases (NVD, CVSS), threat feeds, and proprietary AI enrichment pipelines. This report is provided AS-IS for informational and defensive security purposes only. CYBERDUDEBIVASH accepts no liability for actions taken based on this intelligence. Redistribution, resale, or disclosure to unauthorized parties is strictly prohibited without written consent. For licensing inquiries: bivash@cyberdudebivash.com